
Offline Learning with a Selection Hyper-Heuristic: An Application to Water Distribution Network Optimisation

William B. Yates

wy254@exeter.ac.uk

Computer Science, College of Engineering, Mathematics and Physical Sciences,
University of Exeter, Exeter, EX4 4QF, UK

Edward C. Keedwell

E.C.Keedwell@exeter.ac.uk

Computer Science, College of Engineering, Mathematics and Physical Sciences,
University of Exeter, Exeter, EX4 4QF, UK

https://doi.org/10.1162/evco_a_00277

Abstract

A sequence-based selection hyper-heuristic with online learning is used to optimise 12 water distribution networks of varying sizes. The hyper-heuristic results are compared with those produced by five multiobjective evolutionary algorithms. The comparison demonstrates that the hyper-heuristic is a computationally efficient alternative to a multiobjective evolutionary algorithm. An offline learning algorithm is used to enhance the optimisation performance of the hyper-heuristic. The optimisation results of the offline trained hyper-heuristic are analysed statistically, and a new offline learning methodology is proposed. The new methodology is evaluated, and shown to produce an improvement in performance on each of the 12 networks. Finally, it is demonstrated that offline learning can be usefully transferred from small, computationally inexpensive problems, to larger computationally expensive ones, and that the improvement in optimisation performance is statistically significant, with 99% confidence.

Keywords

Machine learning, selection hyper-heuristics, water distribution networks.

1 Introduction

Optimising the design and rehabilitation of water distribution networks (WDN) is an important real-world problem. A WDN delivers water from reservoirs, tanks, and water treatment facilities to consumers via a network of pipes and makes use of pumps and valves to meet consumer demand. The WDN optimisation problem is characterised as a discrete NP-hard combinatorial optimisation problem with large-scale multimodal search landscapes (see Kheiri et al., 2015).

A variety of metaheuristics (see Blum and Roli, 2003) have been successfully applied to this problem such as simulated annealing (Cunha and Sousa, 1999), shuffled frog leaping algorithms (see Eusuff and Lansey, 2003), harmony search (Geem, 2006), honey-bee mating optimisation (Mohan and Babu, 2010), differential evolution (Zheng et al., 2013), particle swarm optimisation (Ezzeldin et al., 2014), multiobjective evolutionary algorithms (Wang et al., 2015), and ant colony optimisation (Shirzad, 2017). To date, hyper-heuristics have received relatively little attention in the WDN optimisation literature (see Kheiri et al., 2015).

Manuscript received: 8 July 2019; revised: 6 February 2020, 29 May 2020, 2 June 2020, and 11 June 2020; accepted: 11 June 2020.

© 2020 Massachusetts Institute of Technology

Evolutionary Computation 29(2): 187–210

Hyper-heuristics are general purpose heuristic methods. They differ from meta-heuristics in that most metaheuristics search the space of problem solutions, whereas hyper-heuristics search the space of heuristics. In practice this means that a metaheuristic can have access to problem specific information, while a hyper-heuristic is subject to the limitations of the *domain barrier* and is unable to access problem specific information. The domain barrier requires the hyper-heuristic to perform well in the absence of problem specific information and therefore, it is hoped, to be “reuseable” across different problems and problem domains with minimal changes (see Drake et al., 2019).

Hyper-heuristics can be classified as either *generation* or *selection* hyper-heuristics. A generation hyper-heuristic generates new heuristics by discovery, or by modifying or combining existing low-level heuristics. Selection hyper-heuristics, such as those developed in this article, must select and apply a heuristic chosen from a set of low-level heuristics. The goal of both types of hyper-heuristics is to improve the search process through learning and/or optimisation. Such methods have proved effective when applied to a number of real-world problems (see Burke et al., 2019).

Many hyper-heuristics employ learning algorithms in order to improve optimisation performance, and this learning may be classified as either *online* or *offline*. Online learning is based on the low-level heuristic selections and resulting objective function values computed during the execution of a hyper-heuristic on a single problem. In contrast, offline learning is performed on a database of low-level heuristic selections and objective function values computed by a hyper-heuristic on a set of benchmark problems (see Burke et al., 2019).

This article investigates the optimisation of the 12 WDN problems presented in Wang et al. (2015) with the sequence-based selection hyper-heuristic, SSHH described in Kheiri et al. (2015). The SSHH hyper-heuristic employs online learning, and its performance is compared with the results produced by the five multiobjective evolutionary algorithms (or MOEAs) used in Wang et al. (2015). The SSHH hyper-heuristic can also be trained offline using the Baum–Welch learning algorithm detailed in Rabiner (1989) and an appropriate set of heuristic subsequences. The subsequences are chosen from an offline learning database using the statistical methodology presented in Yates and Keedwell (2019). This statistical methodology is also used to analyse the results of offline learning, which leads to an improved offline learning strategy. Finally, the potential for *scalable learning*, where knowledge learned from a small problem is usefully transferred to a second larger problem, is explored.

This study presents four results. Specifically,

1. a hyper-heuristic with online learning is competitive with the state-of-the-art across 12 WDN problems of varying size,
2. offline learning can improve on online learning performance,
3. the effectiveness of the heuristics changes markedly during the optimisation process for WDN problems, and
4. knowledge learned from small problems can be transferred to larger ones, raising the potential for high-performance algorithms trained on benchmarks and deployed on large, real-world problems.

The structure of this article is as follows. Section 2 describes the methodology, Section 3 describes the experimental setup, Section 4 presents the experimental results in detail, and Section 5 contains the conclusions of this study.

Table 1: The size, problem name, acronym, the number of loading conditions (LC), number of water sources (WS), number of decision variables (DV), number of pipe diameter options (PD), for the water distribution network problems. For the TRN problem, three existing pipes have eight diameter options for duplication and the two extra options of cleaning or leaving alone.

Problem	Acronym	LC	WS	DV	PD	Search Space	Size
Two-Reservoir	TRN	3	2	8	8*	3.28×10^7	S
Two-Loop	TLN	1	1	8	14	1.48×10^9	S
BakRyan	BAK	1	1	9	11	2.36×10^9	S
New York	NYT	1	1	21	16	1.93×10^{25}	M
Blacksburg	BLA	1	1	23	14	2.30×10^{26}	M
Hanoi	HAN	1	1	34	6	2.87×10^{26}	M
GoYang	GOY	1	1	30	8	1.24×10^{27}	M
Fossolo	FOS	1	1	58	22	7.25×10^{77}	I
Pescara	PES	1	3	99	13	1.91×10^{110}	I
Modena	MOD	1	4	317	13	1.32×10^{353}	L
Balerna	BIN	1	4	454	10	1.00×10^{455}	L
Exeter	EXN	1	7	567	11	2.95×10^{590}	L

2 Methodology

Section 2.1 contains an overview of the 12 water distribution networks, while Section 2.2 specifies the objective functions to be minimised. Section 2.3 describes the single objective function used in this study, and the solution points that are used to compare optimisation performance. Section 2.4 contains a description of the sequence-based SSHH hyper-heuristic, and the Baum–Welch learning algorithm. Lastly, Section 2.5 presents the elements of the statistical framework introduced in Yates and Keedwell (2019) which are employed in this study.

2.1 Water Distribution Networks

Water distribution networks are an important element of urban infrastructure as they convey clean water from reservoirs, tanks, and water treatment works to homes and businesses via a set of pipes. The design for a WDN aims to ensure a supply of clean water to the demand nodes at sufficient pressure and for minimum monetary cost. Although cost minimisation is the primary design objective, there are many other objectives that can also be considered such as minimisation of water age, adherence to water pressure and velocity constraints, and increasing the robustness of the network to reduce supply outages.

In this article, a simple WDN optimisation problem is considered. The discrete decision variables are the diameters of the pipes in a network. The objectives are to minimise the network’s overall cost and maximise the network’s reliability, while meeting all patterns of demand (or loading) conditions, and maintaining the minimum required head (pressure) throughout the network.

The 12 WDN problems considered here are taken from Wang et al. (2015). Table 1 shows a summary of the problems including the number of loading conditions, water sources, decision variables, and pipe diameter options. The problems are categorised into four groups; small (S), medium (M), intermediate (I), and large (L) according to the




Figure 1: The layout of the Blacksburg network (BLA). The network consists of thirty-five pipes of which twelve have fixed diameters, one reservoir with a fixed head of 715.56 m, and thirty demand nodes. Fixed pipes are shown as blue lines.

size of search space, and range from small benchmark instances with tens of pipes to large-scale city-wide networks comprised of thousands of pipes.

The problems differ from one another in a number of other respects. Of the 12 problems, 11 are based on real-world networks, while the TLN network (see Alperovits and Shamir, 1977) is an example of a hypothetical network. The TRN (see Gessler, 1985), BAK (see Lee and Lee, 2001), NYT (see Schaake and Lai, 1969), and EXN (see Farmani et al., 2004) networks are expansion problems, where the task is to extend an existing network by modifying *some* of the pipes in the network. In addition to selecting pipe diameters such problems can sometimes make use of extra options such pipe cleaning, pipe duplication, or “leaving alone.” The remaining problems TLN, BLA (see Sherali et al., 2001), HAN (see Fujiwara and Khang, 1990), GOY (see Kim et al., 1994), FOS, PES, MOD (see Bragalli et al., 2008), and BIN (see Reca and Martínez, 2006) are pure design problems where the diameters of any or all of the pipes in a network can be modified.

Each problem has minimum head pressure requirements for the demand nodes. The BLA, FOS, PES, and MOD networks also have maximum pressure requirements, and upper bounds on water velocities in the pipes. The TRN network differs from the other problems in that it has three sets of loading conditions, while the BIN network, unlike a typical WDN, has a fixed level of water consumption across all demand nodes.

Figure 1 shows an example schematic of the Blacksburg distribution network.

2.2 Objective Functions

In this study, the water distribution network design problem is specified by three objective functions: the cost C , the head pressure deficit H , and the network’s *resilience*

I_n . The cost and head pressure deficit are to be minimised while the resilience is to be maximised.

The monetary cost is usually expressed in millions and is defined by

$$C = \sum_{i=1}^{np} U_c(D_i)L_i \quad (1)$$

where U_c is the unit pipe cost which depends on the diameter D_i selected, and the length L_i of pipe $i = 1, \dots, np$.

The head pressure deficit is defined by

$$H = \sum_{j=1}^{nn} \left(\max(H_j - H_j^{\max}, 0) + \max(H_j^{\text{req}} - H_j, 0) \right), \quad (2)$$

where H_j is the actual head pressure, H_j^{req} is the minimum required head pressure, and H_j^{\max} is the maximum required head pressure (if any) for each demand node $j = 1, \dots, nn$.

Network resilience measures the redundancy of a pipe network, and maximising this indicator can improve network reliability. It has been shown that using a network resilience index as an additional objective reduces the occurrence of nonviable networks, and yields solutions which are more robust under pipe failure conditions (see Raad et al., 2010). There are however many network resilience measures in the literature, and each has its own particular advantages and disadvantages (see for example Baños et al. (2011)). Each resilience measure attempts to mimic the design goal of designing reliable loops with practicable pipe diameters, while providing excess head pressure above the minimum allowable head pressure at all nodes. In this article, following Wang et al. (2015), a network's resilience is defined by

$$I_n = \frac{\sum_{j=1}^{nn} C_j Q_j (H_j - H_j^{\text{req}})}{\left(\sum_{k=1}^{nr} Q_k H_k + \sum_{i=1}^{npu} \frac{P_i}{\gamma} \right) - \sum_{j=1}^{nn} Q_j H_j^{\text{req}}}, \quad (3)$$

where Q_j is the demand, nr is the number of reservoirs, Q_k is the discharge, and H_k is actual head of reservoir k , npu is the number of pumps, P_i is the power of pump i (if any), and γ is the specific weight of water. The term C_j is the *uniformity* which is defined by

$$C_j = \frac{\sum_{i=1}^{npj} D_i}{npj \max\{D_i\}}, \quad (4)$$

where npj is the number of pipes connected to node j and D_i is the diameter of pipe i connected to node j . The EPANET²¹ software library (Rossman, 2000) is used to run hydraulic simulations in order to obtain the variables required for the calculation of a network's resilience.

2.3 Comparing Solutions

The solutions produced by the SSHH hyper-heuristic are compared with those of the five multiobjective evolutionary algorithms (or MOEAs) used in Wang et al. (2015). It should be emphasised that the objective of this study is not to demonstrate that SSHH

²¹The EPANET software (build 2.00.12) and manual can be downloaded from: <https://www.epa.gov/water-research/epanet/>

Table 2: The minimum cost C (M) and resilience I_n for the viable solutions of the WDN problems found by the MOEAs, and the number of objective function evaluations used to generate them (see Wang et al., 2015).

Prob.	C	I_n	Evals.
TRN	1.7501	0.1490	100,000
TLN	0.4190	0.1535	100,000
BAK	0.9036	0.4978	100,000
NYT	38.8142	0.3906	600,000
BLA	0.1183	0.4267	600,000
HAN	6.1952	0.2041	600,000
GOY	0.1770	0.3262	600,000
FOS	0.0296	0.5239	1,000,000
PES	1.8134	0.2655	1,000,000
MOD	2.5394	0.3619	2,000,000
BIN	1.9986	0.3935	2,000,000
EXN	16.2722	0.3772	2,000,000

is a “superior” optimiser to the MOEAs. Rather, it is to show that SSHH is a computationally efficient optimisation algorithm capable of producing high-quality solutions comparable to the state-of-the-art, and that these solutions can be improved upon with offline learning.

As the SSHH employed in this study is a single objective optimiser the three quantities of cost, head, and resilience are combined to define the single value objective function

$$f = aC + bH + (-cI_n), \tag{5}$$

that is to be minimised. The constants $a = 200$, $b = 1000$, and $c = 5$ are chosen to ensure that the objective function f is always positive across *all* 12 problems, as opposed to “tuning” them for each problem. The objective function favours low-cost networks with low head pressure deficits, as solutions with non-zero deficits are considered to be nonviable.

An MOEA is a multiobjective optimiser that operates on a population of solutions. Such algorithms naturally generate a Pareto front of “best” solutions, which make the trade-offs between the conflicting objectives of cost and resilience explicit. Although SSHH operates on a single solution, it can also generate a Pareto front. However, the use of a single objective value forces the algorithm to explore a smaller region of the solution space; the region of low cost, and therefore lower resilience networks. As a result, comparing the Pareto fronts produced by the two methods is unhelpful. Instead, the SSHH and MOEA optimisers are compared on a single point on the published Pareto front for each problem; the solutions with the cheapest monetary cost for each problem.

Table 2 shows the cheapest viable solutions, their resilience, and the number of objective function evaluations used to generate them, taken from the Pareto fronts of C and I_n presented² in Wang et al. (2015). The Pareto fronts were generated by executing the five MOEAs 30 times, on each problem for the number of iterations shown.

²Implementations of the problems together with their Pareto fronts can be downloaded from: <https://emps.exeter.ac.uk/engineering/research/cws/resources/benchmarks/> under Design/Resilience.

2.4 A Sequence-Based Selection Hyper-Heuristic

A *selection hyper-heuristic* selects heuristics from a given set of low-level heuristics and applies them sequentially to optimise a particular problem. A selection hyper-heuristic can be viewed as an abstraction of an evolutionary algorithm. For example, a genetic algorithm can be represented as a selection hyper-heuristic where the low level heuristics used are crossover and mutate.

The SSHH hyper-heuristic is a sequence-based selection hyper-heuristic with on-line learning (see Kheiri et al., 2015; Kheiri and Keedwell, 2015, 2017). It uses a hidden Markov model (HMM) (see Rabiner, 1989) to generate sequences of heuristic selections, their parameters, and solution acceptance check decisions.

The HMM consists of a set of hidden states, and four probability matrices:

1. a state transition matrix to determine the probability of moving from one hidden state to another,
2. a low-level heuristic emission matrix to determine which low-level heuristic to select,
3. a parameter emission matrix to determine the parameter for a low-level heuristic, and
4. an acceptance check emission matrix to determine when a solution should be evaluated and checked for acceptance.

In the absence of a priori knowledge regarding a particular problem the number of hidden states is set to be the number of low-level heuristics in the domain, and the state transition, the parameter and acceptance check matrices are set to be equiprobable. The low-level heuristic emission matrix is set to the identity matrix. This ensures that, initially, each equiprobable hidden state emits a single low-level heuristic together with an equiprobable choice of heuristic parameter and acceptance check decision.

At each iteration of the optimisation process, the SSHH hyper-heuristic moves from its current state to a new state according to the probabilities of the transition matrix. Once a new state has been chosen, the emission probability matrices are used to determine a low-level heuristic, its parameters, and whether to check for acceptance or not. The selected parametrised low-level heuristic is applied to the current solution in order to derive a new solution. If the acceptance check decision is true, the objective function is evaluated on the new solution, and a decision is made whether to accept it. Specifically, a new solution is accepted if it improves on the current solution or is within 5% of the current best solution; otherwise it is rejected.

The SSHH hyper-heuristic employs *online* learning. During optimisation, SSHH keeps a history of the heuristic selections, parameters, and acceptance checks produced by the HMM. If, following an acceptance check, a new, best solution is found, the online learning algorithm steps through the history, increasing the probabilities of the accepted state transitions and emissions that led to the new minima. Thus the probability that the SSHH produces the sequence of heuristic selections and emissions contained in the history is now higher. After the acceptance check, the history is erased and the optimisation process is resumed.

The SSHH hyper-heuristic can also be trained *offline* using the Baum–Welch learning algorithm (see Rabiner, 1989). The Baum–Welch algorithm calculates the maximum likelihood estimates of a HMM’s parameters for an arbitrary sequence of observations. In this study, the parameter to be estimated is the state transition probability matrix, and

the observations are sequences of heuristics chosen from the offline learning database. The objective is to demonstrate that the Baum–Welch algorithm can learn an effective optimisation strategy for a problem before online learning refines the approach.

The SSHH hyper-heuristic requires a number of low-level heuristics. The five heuristics used in Kheiri et al. (2015) and a two-point crossover heuristic C_5 are used in these experiments. The low-level heuristics are:

1. M_0 – change one pipe diameter randomly,
2. S_1 – swap two pipe diameters at random,
3. M_2 – increase or decrease a randomly selected pipe diameter by one size,
4. R_3 – “ruin” several pipes and rebuild randomly where the number of pipes to be changed is a parameter in the range $[1,5]$,
5. S_4 – shuffle several pipes (i.e., makes several swaps) where the number of pipes to be changed is a parameter in the range $[1,5]$, and
6. C_5 – two-point crossover of two vectors of network pipe diameters.

2.5 Logarithmic Returns

Although each WDN problem has the same objective function f to be minimised, the range of f varies considerably between problem instances. Without a priori knowledge of the objective function ranges, the objective function values from different problems cannot be compared directly. Instead, following the methodology introduced in Yates and Keedwell (2019), normalised subsequences of objective function values are compared. The definitions of the functions $\bar{\alpha}$ and $\bar{\alpha}_f$ employed in this study are reproduced here.

Consider a series of objective function values o_0, o_1, \dots, o_n observed after applying a subsequence s of n low-level heuristics to some initial solution x_0 . The *log return* α of subsequence s is

$$\alpha(s) = \log_{10} \left(\frac{o_n}{o_0} \right). \quad (6)$$

The *mean log return* of a set of N subsequences is defined by

$$\bar{\alpha}(\{s_1, \dots, s_N\}) = \frac{1}{N} \sum_{i=1}^N \alpha(s_i). \quad (7)$$

The *final log return* of a hyper-heuristic experiment or run is the log return between the initial solution x_0 and the final best solution x_{\min} found during the run. In symbols

$$\alpha_f(s) = \log_{10} \left(\frac{o_{\min}}{o_0} \right), \quad (8)$$

where o_{\min} is the objective function value of x_{\min} . The overall performance of a hyper-heuristic is measured by the *mean final log return* $\bar{\alpha}_f$ of a set of N sequences defined by

$$\bar{\alpha}_f(\{s_1, \dots, s_N\}) = \frac{1}{N} \sum_{i=1}^N \alpha_f(s_i). \quad (9)$$

The functions $\bar{\alpha}$ and $\bar{\alpha}_f$ are the means of log values. The anti-log of the mean of the logs is equivalent to the *geometric mean*. The geometric mean is used so that no range

Algorithm 1 The DBGen hyper-heuristic in pseudocode

```

1.  $ITERATIONS \leftarrow MAX\_ITER$ 
2.  $new\_sol \leftarrow initialiseSolution()$ 
3.  $new\_obj \leftarrow f(new\_sol)$ 
4. for  $i = 1$  to  $5$  do
5.    $cross\_sol[i] \leftarrow initialiseSolution()$ 
6.    $cross\_obj[i] \leftarrow f(cross\_sol[i])$ 
7. end
8. while ( $ITERATIONS > 0$ ) do
9.    $cur\_sol \leftarrow new\_sol$ 
10.   $cur\_obj \leftarrow new\_obj$ 
11.  Heuristic  $h \leftarrow selectHeuristic()$ 
12.   $i \leftarrow ranInt()$ 
13.   $new\_sol \leftarrow apply(h, new\_sol, cross\_sol[i])$ 
14.   $new\_obj \leftarrow f(new\_sol)$ 
15.  if ( $new\_obj \geq cur\_obj$  and  $ranFloat() \geq 0.5$ ) then /* reject new solution */
16.     $new\_sol \leftarrow cur\_sol$ 
17.     $new\_obj \leftarrow cur\_obj$ 
18.  else
19.     $i \leftarrow ranInt()$ 
20.     $cross\_sol[i] \leftarrow new\_sol$ 
21.     $cross\_obj[i] \leftarrow new\_obj$ 
22.  end
23.   $ITERATIONS \leftarrow ITERATIONS - 1$ 
24. end

```

dominates the average, and for this reason, in this study, $\bar{\alpha}$ and $\bar{\alpha}_f$ are used in preference to the arithmetic mean of the decimal returns when comparing results across multiple problems.

3 Experimental Setup

Section 3.1 introduces the DBGen hyper-heuristic that is used to generate the offline learning database, while Section 3.2 contains the definition of the γ -ratio which is used to select training subsequences from the database. Section 3.3 details the training and testing methodology used to demonstrate generalisation across the WDN problems. Lastly, in Section 3.4, the statistical test used to validate improvements in the optimisation results is described.

All experiments were conducted on a Mac Pro computer with a 3.5-GHz, 6-core, Intel Xeon E5 processor, and 16 GB of 1866 MHz memory.

3.1 An Offline Subsequence Database

The unbiased, random, single selection hyper-heuristic DBGen used to generate the database of heuristic selections and objective function values is shown in listing 1.

The function *ranInt()* (lines 12 and 19) returns a uniformly distributed pseudorandom number in the set $\{1, 2, 3, 4, 5\}$, while the function *ranFloat()* (line 15) returns a uniformly distributed pseudorandom number in the interval $(0,1)$. The function *selectHeuristic()* (line 11) selects a single heuristic class at random from the set $\{M, S, R, C\}$. The function *apply()* (line 13) takes the heuristic class and chooses, again at random, a low-level heuristic and its parameters, from the available heuristics of that class. The low-level heuristic is then applied to the current solution, and if the class is *C*, to a crossover solution which is randomly selected from the crossover pool (lines 5–6 and lines 12–13). An objective function evaluation (line 14) and an acceptance check (line 15) are then performed. If a new solution's objective value is less than the current solution's objective value or *ranFloat()* < 0.5 then it is accepted, and also stored, at random, in the

crossover pool (lines 19–21). Otherwise the new solution is rejected. The random term allows new solutions to be accepted regardless of their objective function for 50% of the cases. Accepting states that may lead to a large increase in objective function value forces the DBGen hyper-heuristic to explore the space of low-level heuristic selections instead of optimising the problem efficiently.

When crossover heuristics are available, the choice of crossover mechanism also affects hyper-heuristic performance (see Drake et al., 2015). The DBGen crossover mechanism (and the number of crossover solutions) is taken from the crossover management scheme employed by the AdapHH hyper-heuristic (see Drake et al., 2015). This crossover mechanism is also used by SSHH.

The DBGen hyper-heuristic is executed 40 times on each of the 12 WDN problem instances for 10,000, 20,000, 50,000, and 100,000 iterations for the small, medium, intermediate, and large problems, respectively. The number of iterations (denoted by MAX_ITER in listing 1) for each size were chosen for computational feasibility. The number of 40 runs was chosen so as to ensure that robust statistics could be calculated for each problem instance.

For each problem, each DBGen run is seeded by a unique number which is used to initialise the pseudorandom number generator used by DBGen, and to randomly initialise a WDN problem instance.

Each run or sequence is then broken down into consecutive subsequences of length $n = 2, 3, \dots, 10$. For example, given a sequence $\{M_0S_4R_3M_2S_1\}$ of length five, the sets of subsequences of lengths two, three, and four are

$$\{M_0S_4, S_4R_3, R_3M_2, M_2S_1\}, \{M_0S_4R_3, S_4R_3M_2, R_3M_2S_1\}, \text{ and } \{M_0S_4R_3M_2, S_4R_3M_2S_1\},$$

respectively.

Although the computational cost of constructing such a database is significant, the objective of this study is not to compare optimisation performance using equivalent computational resources, but rather (as mentioned in Section 2.2) it is to investigate what can be learned from a preexisting offline learning database that has been constructed using a method that is independent of the optimisation algorithm under consideration.

3.2 Selecting Subsequences with the γ -Ratio

The function used to select effective subsequences from the offline learning database, which is introduced in Yates and Keedwell (2019), is defined here.

The *unit log return* of a set of N subsequences s_i of length n_i is

$$\beta(\{s_1, \dots, s_N\}) = \sum_{i=1}^N \frac{\alpha(s_i)}{n_i}. \quad (10)$$

The length of each subsequence is important because for some problems the execution times of the low-level heuristics and objective function evaluations can be non-trivial.

Let S^+ be the set of all subsequences with $\alpha(s) > 0$, and let S^- be the set of all subsequences with $\alpha(s) \leq 0$, and note that the set of all subsequences $S = S^+ \cup S^-$. The positive and negative parts of β can be separated out by $\beta = \beta^+ - \beta^-$ where

$$\beta^+(U) = \beta(U \cap S^+) \text{ and } \beta^-(U) = -\beta(U \cap S^-), \quad (11)$$

for every subset U of S .

A subsequence s may occur a number of times in the database and each instance or occurrence will have a different subsequence of objective function values depending

on the problem, the run, and the position in a run where s arises. The set

$$U_s = \{s^1, \dots, s^{N_s}\}, \quad (12)$$

is the set of all occurrences of s , where N_s is the number of occurrences.

In this study the γ -ratio

$$\gamma(s) = \frac{\beta^-(U_s)}{\beta^+(U_s) + 1}, \quad (13)$$

is used to select subsequences from the offline database. Large values of $\gamma(s) > 1$ indicate an effective subsequence which tends to decrease the objective function value, while small values of $\gamma(s) < 1$ indicate a disruptive subsequence, which tends to increase the objective function value.

3.3 Offline Learning and Generalisation

The Baum–Welch learning algorithm is used to estimate the parameters of a HMM in order to maximise the probability of generating a given sequence of training observations. In this study, the parameter to be estimated is the state transition probability matrix of the SSHH hyper-heuristic, and the training observations are subsequences of low-level heuristics that are selected from the offline learning database using the γ -ratio. Specifically, following the method presented in Yates and Keedwell (2019), the 10 subsequences of lengths two and three with the largest γ -ratio are chosen. Subsequences of lengths two and three are used as they occur more frequently in the offline learning database than longer subsequences, and as a result, the statistics calculated for these subsequences are more reliable. The number of 10 subsequences was chosen to ensure that the subsequence set contained sufficient heuristic “diversity,” that is, the subsequences consist of more than just one or two low-level heuristics.

The results of executing the offline trained SSHH hyper-heuristic on the WDN problems are compared with those of an untrained SSHH hyper-heuristic using a *leave-one-out* cross-validation methodology (see Bishop, 2006). Recall that there are 12 problem instances in WDN domain. For each target problem, the training set consists of the 10 subsequences with the largest γ -ratios, chosen from the subsequences of the remaining 11 problems. The subsequences are used to train the HMM of the SSHH hyper-heuristic which is then evaluated on the target problem. This ensures that SSHH is always evaluated on a problem that is “unseen.” The objective is to demonstrate empirically that an offline trained SSHH hyper-heuristic is able to learn and generalise from training subsequences selected across a number of problems. In this context, generalisation means that the trained SSHH hyper-heuristic is able to significantly outperform the untrained SSHH hyper-heuristic when evaluated on unseen problem instances.

After Baum–Welch training, the probability matrices are edited to ensure that every state transition and heuristic emission has a non-zero probability of at least 0.0001. The trained SSHH hyper-heuristic, denoted T-SSHH, is initialised with these edited matrices.

3.4 Assessing Hyper-Heuristic Performance

In the evolutionary computation literature statistical tests are widely used to compare and rank the performance of algorithms that have been evaluated on a number of benchmark problems (see, for example, Dietterich, 1998; Demšar, 2006; Derrac et al., 2011; and Veek et al., 2017).

In this article, following Demšar (2006), the non-parametric, one-tailed Wilcoxon signed-rank test is used to establish a *stochastic ordering* on two hyper-heuristics A and

B. The null hypotheses of the Wilcoxon test is that the median difference between pairs of observations is zero. The null hypothesis is tested at a significance level of 0.01 on sample sizes of 480, and is rejected if the p -value is less than 0.01. In this case, the alternative hypothesis that the median difference between pairs of observations is *less* than zero is accepted with 99% confidence. This implies that the random variable $\bar{\alpha}_f(A)$ is “smaller” than the random variable $\bar{\alpha}_f(B)$, and thus hyper-heuristic A is more effective than hyper-heuristic B. The $\alpha_f(A)$ and $\alpha_f(B)$ values can be paired because the initial seed and therefore the initial solution for each problem instance $p = 1, \dots, 12$, and each run $r = 1, \dots, 40$ is the same for both hyper-heuristics.

A more detailed discussion of the Wilcoxon test and its appropriateness for this work can be found in the Appendix.

4 Results

In Section 4.1, the SSHH hyper-heuristic is used to optimise the 12 WDN problems. In Section 4.2, in order to improve optimisation performance the SSHH hyper-heuristic’s HMM is trained offline with the Baum–Welch learning algorithm. However, in this case, offline learning fails to improve optimisation performance. With this in mind, Section 4.3 presents an analysis of the results of offline learning, which leads, in Section 4.4, to an effective offline learning methodology. Finally, in Section 4.5, the potential for scalable learning is explored.

4.1 Online Learning

In this section, the SSHH hyper-heuristic is used to optimise the 12 WDN problems. This experiment extends the work presented in Kheiri et al. (2015) by evaluating the SSHH hyper-heuristic, with an additional crossover operator, on multiple WDN problems. The objective is to compare the performance of the SSHH algorithm with that of the evolutionary algorithms, and to determine the utility of the extra crossover heuristic.

The SSHH hyper-heuristic is executed 40 times on each of the 12 problems in the WDN domain. For each problem, each run is seeded by a number that is distinct to the seeds used to generate the offline learning database. The number of iterations used by SSHH varies with the problem size. Specifically, SSHH is executed for 10,000, 20,000, 50,000, and 100,000 iterations for the small, medium, intermediate, and large problems, respectively. The number of iterations for each size were chosen for computational feasibility. The number of 40 runs was chosen so as to ensure that robust statistics could be calculated for each problem instance. The SSHH hyper-heuristic is compared to the MOEAs on a single point on the published Pareto fronts; the solutions with the cheapest monetary cost. The cheapest solutions found by each algorithm for each problem are shown in Table 3. As the SSHH only evaluates the objective function after an acceptance check, rather than every iteration, the number of objective function evaluations (Evals.) for SSHH is always less than the number of iterations.

The cheapest solutions found by SSHH and the MOEAs for the problems BAK and NYT are identical. For the small problem TRN and TLN, the medium-sized problems BLA, HAN, GOY, and FOS no solution dominates. For the intermediate problem PES, and the large problems MOD and BIN, the MOEA solutions (shown in boldface) dominate those found by SSHH.

The solution produced by SSHH for EXN has a head deficit of $H = 2.6940$. The head penalties are calculated by summing the pressure violations over the whole network. When the violations are small, and spread evenly across a network, the solution can be viewed as semi-viable, or approaching viability. For the EXN network, the pressure

Table 3: The lowest cost C (M), resilience I_n , and the number of objective function evaluations (Evals.) for the solutions of the WDN problems produced by SSHH and MOEA. The result R indicates an equal (E), non-dominant (ND), dominant (D), or nonviable (NV) solution.

Prob.	SSHH			MOEAs			R
	C	I_n	Evals.	C	I_n	Evals.	
TRN	1.7501	0.1110	8876	1.7501	0.1490	100000	ND
TLN	0.4200	0.1579	5850	0.4190	0.1535	100000	ND
BAK	0.9036	0.4978	9257	0.9036	0.4978	100000	E
NYT	38.8142	0.3906	18282	38.8142	0.3906	600000	E
BLA	0.1186	0.4804	17375	0.1183	0.4267	600000	ND
HAN	6.1350	0.1797	18580	6.1952	0.2041	600000	ND
GOY	0.1781	0.4498	17334	0.1770	0.3262	600000	ND
FOS	0.0296	0.5249	49124	0.0296	0.5239	1000000	ND
PES	1.8319	0.2210	48438	1.8134	0.2655	1000000	D
MOD	2.5754	0.2739	81101	2.5394	0.3619	2000000	D
BIN	2.1366	0.3280	56005	1.9986	0.3935	2000000	D
EXN	7.6130	0.1912	67841	16.2722	0.3772	2000000	NV

Table 4: The lowest cost C (M), head deficit H , and resilience I_n for the solutions of the PES, MOD, BIN, and EXN* problems produced by SSHH and MOEA. The result R indicates a non-dominant (ND) or dominant (D) solution.

Prob.	SSHH			MOEA			R
	C	I_n	Evals.	C	I_n	Evals.	
PES	1.8125	0.2546	557549	1.8134	0.2655	1000000	ND
MOD	2.5485	0.2792	1457855	2.5394	0.3619	2000000	D
BIN	2.0919	0.3440	1014035	1.9986	0.3935	2000000	D
EXN*	15.5632	0.3425	681349	16.2722	0.3772	2000000	ND

violation occurs at a single pipe, and so the SSHH solution is deemed nonviable. However, by choosing the alternative constants $a = 2$, $b = 5000$, and $c = 0.1$ for the objective function, SSHH can find a viable solution where $C = 17.0886$ and $I_n = 0.3373$, using 49552 objective function evaluations. This result is denoted EXN*.

Table 4 shows the cheapest cost solutions found by executing SSHH on PES, MOD, BIN, and EXN for additional iterations. Specifically, SSHH is run 40 times for 1,000,000 iterations on PES and 2,000,000 iterations on MOD, BIN and EXN. As extending the number of iterations for EXN using the original objective function constants also yields a nonviable solution, the result shown in Table 4, denoted EXN*, uses the alternative constants defined above. While the MOD and BIN solutions remain dominant after the additional iterations, the solutions for PES and EXN* are now non-dominated. For PES, SSHH actually finds a lower cost solution than any of the MOEAs.

As the constants used to define the objective function are chosen to ensure that the function is positive for *all* the problems in the WDN domain, one would expect a

Table 5: A problem-by-problem comparison of the mean final objective function value and standard deviation for SSHH and T-SSHH. Winning scores are shown in boldface.

Prob.	SSHH	SD	T-SSHH	SD
TRN	357.4068	16.9400	361.0980	18.6944
TLN	87.2512	2.6345	87.9132	2.4008
BAK	179.0267	0.9965	179.2223	1.4257
NYT	8116.0158	536.2183	8169.6525	645.3373
BLA	22.2911	1.2971	22.5268	2.4840
HAN	1286.6305	37.4569	1294.9970	42.7782
GOY	33.3999	0.1606	33.3900	0.0628
FOS	4.2449	0.6641	4.4674	0.9202
PES	401.5132	31.3177	409.5692	50.8569
MOD	555.0569	27.6637	564.3754	100.4345
BIN	494.2631	65.7123	537.6853	115.3677
EXN	4789.5943	402.1191	4738.0677	810.7313

multiobjective algorithm to outperform a single objective function optimiser. However, the differences in cost and network resilience are modest, while SSHH employs significantly less objective function evaluations than the MOEAs. This is important because, for many large problems (such as EXN), evaluating the objective function is computationally expensive.

These results demonstrate that the SSHH hyper-heuristic is a computationally efficient alternative to a multiobjective evolutionary algorithm for the optimisation of WDN problems. Furthermore the SSHH algorithm can be extended to optimise multiple objectives (see Walker and Keedwell, 2016).

4.2 Offline Learning

In order to improve optimisation performance the SSHH hyper-heuristic’s HMM is trained offline with the Baum–Welch learning algorithm. The training observations consist of effective subsequences of low-level heuristics selected from the offline database using the γ -ratio.

The SSHH hyper-heuristic is trained offline with the Baum–Welch learning algorithm on the 10 subsequences of heuristics of lengths two and three with the largest γ -ratios in the offline learning database. The training sets are constructed from these effective subsequences using a leave-one-out cross-validation methodology. Specifically, for each WDN problem, the training set consists of the 10 subsequences with the largest γ -ratios, chosen from the subsequences of the remaining problems. These subsequences are used to offline train the HMM of the SSHH hyper-heuristic which is then used to optimise the “unseen” target problem. This methodology gives rise to 12 training sets, one for each of the WDN problems.

The offline trained hyper-heuristic T-SSHH is executed 40 times on each of the 12 problems in the WDN domain. The mean final objective function values for each problem are shown in Table 5.

The problem-by-problem results demonstrate that the SSHH hyper-heuristic outperforms the offline trained hyper-heuristic T-SSHH on 10 of the 12 WDN problems. The overall results are shown in Table 6.

Table 6: The mean final log return $\bar{\alpha}_f$, the mean percentage change, the mean number of iterations to a minimum, and the mean number of objective function evaluations.

	$\bar{\alpha}_f$	%	Min. Sel.	Obj. Eval.
SSHH	-2.9104	-88.7454	32916.5854	32238.0979
T-SSHH	-2.9134	-88.7435	32695.6125	30071.8729

Table 7: The sample pseudo-median difference \hat{d} , the sample median absolute deviation MAD, the sample mean difference \bar{d} , the standard deviation SD, the p -value, and the 99% confidence interval for $\bar{\alpha}_f(\text{T-SSHH}) - \bar{\alpha}_f(\text{SSHH})$. Statistically significant results are shown in boldface.

\hat{d}	MAD	\bar{d}	SD	p -value	Conf. Int.
0.0017	0.0225	-0.0030	0.2011	0.9029	$[-\infty, 0.0055]$

Although the mean final log return $\bar{\alpha}_f$ is slightly better for T-SSHH than SSHH, the mean percentage change in optimisation value is worse. The differences in final log returns are tested for statistical significance as explained in Section 3.4 and the results are shown in Table 7.

The results of the one-tailed Wilcoxon test together with the results contained in Tables 5 and 6 indicate that T-SSHH is an inferior optimiser when compared with SSHH, and that offline learning has failed to improve optimisation performance.

4.3 An Analysis of Heuristic Effectiveness

In Yates and Keedwell (2019) the γ -ratio is used to select effective subsequences of heuristics from a number of distinct problem domains. The γ -ratio can also be used, in certain circumstances, to select effective subsequences *across* a number of problem domains. For example, consider two comparable domains, and define the $\bar{\alpha}$ -order of a domain to be the order of its low-level heuristics when ranked by their mean log return $\bar{\alpha}$. If the difference between the $\bar{\alpha}$ -orderings of the two domains is small, then a subsequence that is effective in one domain is likely to be effective in the other. In this section, the failure of offline training to improve optimisation performance in the previous section is investigated by considering the $\bar{\alpha}$ -order of the low-level heuristics of the WDN domain *during* the course of the optimisation process.

Two sets of heuristic instances are selected from the offline learning database. These sets contain the heuristic instances that occur at the “beginning” of the optimisation process, when objective function values are relatively high, and at the “end” of the optimisation process, when objective function values are relatively low.

The *current objective function* value for a heuristic (or subsequence of heuristics) in a run is the objective function value o_t at time t prior to applying the heuristic (subsequence). The sets LOW and HIGH consist of all those heuristic instances which have a current objective function value

$$o_t < P_{10}^i \text{ and } o_t > P_{90}^i, \tag{14}$$

respectively, where P_{10}^i is the 10th and P_{90}^i is the 90th percentile. The set LOW contains the 10% of heuristic instances with the lowest current objective function values while

Table 8: The low-level heuristics in the HIGH and LOW subsets ordered by ascending mean log return $\bar{\alpha}$ from left to right, the Spearman's Footrule distance d , and the normalised Footrule distance $d' = \frac{d}{m}$.

Prob.	Set	$\bar{\alpha}$ -order	d	d'
TRN	HIGH	C ₅ , R ₃ , S ₄ , M ₀ , S ₁ , M ₂	10	0.5556
	LOW	C ₅ , M ₂ , M ₀ , R ₃ , S ₁ , S ₄		
NYT	HIGH	C ₅ , S ₄ , R ₃ , S ₁ , M ₀ , M ₂	12	0.6667
	LOW	C ₅ , M ₂ , M ₀ , R ₃ , S ₁ , S ₄		
FOS	HIGH	S ₄ , C ₅ , R ₃ , S ₁ , M ₀ , M ₂	16	0.8889
	LOW	M ₂ , M ₀ , C ₅ , S ₁ , R ₃ , S ₄		
MOD	HIGH	C ₅ , R ₃ , M ₀ , M ₂ , S ₁ , S ₄	6	0.3333
	LOW	M ₂ , C ₅ , M ₀ , R ₃ , S ₁ , S ₄		
EXN	HIGH	C ₅ , R ₃ , M ₀ , M ₂ , S ₁ , S ₄	6	0.3333
	LOW	C ₅ , M ₂ , M ₀ , S ₁ , R ₃ , S ₄		
ALL	HIGH	C ₅ , S ₄ , R ₃ , S ₁ , M ₀ , M ₂	14	0.7778
	LOW	M ₂ , C ₅ , M ₀ , S ₁ , R ₃ , S ₄		

the set HIGH contains the 10% of heuristic instances with the highest current objective function values.

Calculating the percentile values over all 480 sequences in the offline learning database can lead to some problem instances dominating the sets; those that produce very high or very low objective function values. To mitigate this, the percentiles P^i are calculated locally over the objective function values of *each* sequence $i = 1, \dots, 480$. This ensures that heuristic instances from all the problems in the domain are included in the LOW and HIGH sets.

The mean log return $\bar{\alpha}$ of the low-level heuristics is calculated from the LOW and HIGH sets. The heuristics are then ranked by their $\bar{\alpha}$. The resulting LOW and HIGH orderings for the smallest problem in each size category, EXN, and overall, are shown in Table 8.

The change in orderings can be quantified by using the Spearman's Footrule distance (see Diaconis and Graham, 1977). The distance is calculated by taking the sum of the absolute values of the difference between two ranks. In symbols, if σ and π denote two permutations of n elements such that $\sigma(i)$ and $\pi(i)$ denote the rank of an element $i = 1, \dots, n$ in the permutation, then Spearman's Footrule is defined by

$$d(\sigma, \pi) = \sum_{i=1}^n |\sigma(i) - \pi(i)|, \quad (15)$$

and has a maximum integer value of $m = \lfloor \frac{1}{2}n^2 \rfloor$ where $\lfloor \cdot \rfloor$ is the floor function. The greater the Footrule distance, the greater the difference between the two orders.

The results in Table 8 show some large differences in the orderings of the LOW and HIGH heuristic sets. For example, notice how the M₂ heuristic changes from being one of the least effective heuristics in the HIGH sets, to one of the most effective heuristics in the LOW sets. These large differences in rank indicate that different individual heuristics are effective at different points in the optimisation process. Figure 2 shows the effectiveness of the low-level heuristics for each percentile over all 12 problems. The plot illustrates the changes in heuristic effectiveness as solution optimality increases. Notice




Figure 2: The mean log return $\bar{\alpha}$ of the low-level heuristics in the WDN domain, averaged over the 10 local percentiles. Optimality increases from left to right, while negative $\bar{\alpha}$ values correspond to reductions in the objective function value.

that the two-point crossover heuristic C_5 is the best performing low-level heuristic in all but the last two percentiles: P_{20}^i and P_{10}^i .

This large change in heuristic effectiveness is particularly relevant to SSHH as the efficient optimisation of such problems require online learning strategies, as any offline learned heuristic subsequences can only be effective at particular points during optimisation. The result of this analysis is also important for other optimisation techniques. For example, a vanilla genetic algorithm will typically execute a mutation operation and crossover operation at a given fixed rate for every iteration of the algorithm. For problems where the effectiveness of the crossover and mutate operation varies significantly during optimisation, optimisation performance can be improved by varying this rate accordingly.

4.4 Offline Learning with LOW Subsequences

In order to test the hypothesis that it is the large variance in heuristic performance during optimisation that prevents the Baum–Welch algorithm from finding a suitable set of HMM parameters, the experiment in Section 4.2 is repeated using subsequences that are effective when the objective function value is relatively low. Specifically, the 10 subsequences of lengths two and three with the largest γ -ratios are selected from the LOW set defined in Section 4.3. The effective LOW subsequences are then used by the Baum–Welch algorithm to train a HMM using a leave-one-out cross-validation methodology.

The SSHH hyper-heuristic is initialised with an identity heuristic emission matrix and equiprobable transition, parameter and acceptance matrices. The optimisation process with online learning is started, and when half of the specified iterations have been performed, the current HMM is switched for the offline trained HMM. The optimisation process, again with online learning, is then resumed. Although the method of switching to the LOW trained HMM is simple, it is trivial to implement, and requires no information regarding the objective function. The results for the SSHH hyper-heuristic when

Table 9: A problem-by-problem comparison of the mean final objective function value and standard deviation for SSHH and T-SSHH-L. Winning scores are shown in boldface.

Prob.	SSHH	SD	T-SSHH-L	SD
TRN	357.4068	16.9400	355.4825	14.4080
TLN	87.2512	2.6345	86.7689	2.4960
BAK	179.0267	0.9965	178.8902	0.9436
NYT	8116.0158	536.2183	7985.0690	448.5106
BLA	22.2911	1.2971	22.0291	1.1858
HAN	1286.6305	37.4569	1279.6965	36.1578
GOY	33.3999	0.1606	33.3729	0.0081
FOS	4.2449	0.6641	4.1938	0.6468
PES	401.5132	31.3177	393.3359	27.0505
MOD	555.0569	27.6637	545.1349	24.2674
BIN	494.2631	65.7123	474.7110	34.7078
EXN	4789.5943	402.1191	4667.1560	272.0995

Table 10: The mean final log return $\bar{\alpha}_f$, the mean percentage change, the mean number of iterations to a minimum, and the mean number of objective function evaluations.

	$\bar{\alpha}_f$	%	Min. Sel.	Obj. Eval.
SSHH	-2.9104	-88.7454	32916.5854	32238.0979
T-SSHH-L	-2.9248	-88.8740	34693.5292	30437.8688
T-SSHH-L-NS	-2.9105	-88.7656	32384.0750	32212.7813

it is trained offline using effective subsequences chosen from the LOW set are denoted T-SSHH-L. The results for the offline trained hyper-heuristic with *no switching* mechanism, where the offline trained HMM is used from the outset of the optimisation process, are included for comparison purposes, and are denoted T-SSHH-L-NS.

The offline trained hyper-heuristic T-SSHH-L and T-SSHH-L-NS are each executed 40 times on each of the 12 problems in the WDN domain. Table 9 contains the mean final objective function value for each problem in the WDN domain. The T-SSHH-L hyper-heuristic outperforms SSHH on all 12 problems.

Overall, the results in Table 10 show that T-SSHH-L outperforms SSHH. The T-SSHH-L hyper-heuristic also outperforms T-SSHH-L-NS which demonstrates the importance of applying the effective subsequences at an appropriate point in the optimisation process.

The overall differences in the final log returns are tested for statistical significance as before, and the results are shown in Table 11. They demonstrate that the differences, while small, are statistically significant, and so the offline trained hyper-heuristic T-SSHH-L outperforms the SSHH hyper-heuristic with 99% confidence.

These results show that, despite large variations in heuristic performance over the optimisation process, it is possible to significantly improve the optimisation performance of the SSHH hyper-heuristic on unseen WDN problems with offline learning. The result for EXN is particularly encouraging, as it shows that a training set constructed

Table 11: The sample median difference \hat{d} , the sample median absolute deviation MAD, the sample mean difference \bar{d} , the standard deviation SD, the p -value, and the 99% confidence interval for $\bar{\alpha}_f(\text{T-SSHH-L}) - \bar{\alpha}_f(\text{SSHH})$. Statistically significant results are shown in boldface.

\hat{d}	MAD	\bar{d}	SD	p -value	Conf. Int.
-0.0056	0.0015	-0.0144	0.1941	0.0000	$[-\infty, -0.0036]$

Table 12: The lowest cost C (M) and resilience I_n for the solutions of the WDN problems produced by SSHH, T-SSHH-L, and the MOEAs. The solutions for EXN produced by SSHH and T-SSHH-L are nonviable with a head deficit H of 2.6940 and 7.4360, respectively.

Prob.	SSHH		T-SSHH-L		MOEAs	
	C	I_n	C	I_n	C	I_n
TRN	1.7501	0.1110	1.7501	0.1110	1.7501	0.1490
TLN	0.4200	0.1579	0.4200	0.1579	0.4190	0.1535
BAK	0.9036	0.4978	0.9036	0.4978	0.9036	0.4978
NYT	38.8142	0.3906	38.8142	0.3906	38.8142	0.3906
BLA	0.1186	0.4804	0.1186	0.4804	0.1183	0.4267
HAN	6.1350	0.1797	6.1177	0.1764	6.1952	0.2041
GOY	0.1781	0.4498	0.1783	0.4607	0.1770	0.3262
FOS	0.0296	0.5249	0.0296	0.5249	0.0296	0.5239
PES	1.8319	0.2210	1.8276	0.2171	1.8134	0.2655
MOD	2.5754	0.2739	2.5980	0.2775	2.5394	0.3619
BIN	2.1366	0.3280	2.1668	0.2967	1.9986	0.3935
EXN	7.6130	0.1912	2.7001	0.1891	16.2722	0.3772

from a number of smaller problems can lead to improved optimisation on a larger, more complex, problem.

Table 12 contains a comparison of the cheapest monetary cost solutions found by SSHH, T-SSHH-L, and the MOEAs for each problem. Dominant solutions are shown in boldface. The performance of T-SSHH-L is markedly inferior to the MOEAs on the large problems MOD, BIN and EXN. This could be due to the relatively low number of iterations employed by SSHH on these problems. Another cause could be that the 100000 iterations employed by DBGen is insufficient to discover the subsequences of heuristics necessary to construct effective training sets for the large problems.

It should be noted that although the values of C , H and I_n vary considerably across the WDN problems, the objective function constants a , b , and c are identical for each task. The EXN* result (see Table 4) demonstrates that optimisation performance can be improved by “tuning” the constants for a particular problem. However, preliminary experiments suggest that using the same objective function for each problem facilitates learning and generalisation.

Table 13: A problem-by-problem comparison of the mean final objective function value and standard deviation for T-SSHH-L and T-SSHH-L-TRN. Winning scores are shown in boldface.

Prob.	T-SSHH-L	SD	T-SSHH-L-TRN	SD
TRN	355.4825	14.4080	355.6056	13.5715
TLN	86.7689	2.4960	87.0160	2.5482
BAK	178.8902	0.9436	178.8715	0.9204
NYT	7985.0690	448.5106	8011.2947	444.8608
BLA	22.0291	1.1858	22.0194	1.0362
HAN	1279.6965	36.1578	1276.8813	27.1250
GOY	33.3729	0.0081	33.3721	0.0079
FOS	4.1938	0.6468	4.2177	0.6845
PES	393.3359	27.0505	396.5199	29.9579
MOD	545.1349	24.2674	549.7243	25.2320
BIN	474.7110	34.7078	487.6116	39.1680
EXN	4667.1560	272.0995	4632.3642	333.9616

4.5 Scalable Learning

In this section, the potential for *scalable learning* is explored. Scalable learning is where a model developed for a small, computationally tractable task, is reused as the starting point of a model for a larger, second task. The concept of scalable learning is appropriated from Burke et al. (2007) where the authors use evolutionary algorithms to generate novel heuristics for small problems which are then shown to perform well when tested on larger problems.

The idea is to generate training subsequences from a small problem which are then used, offline, to improve optimisation performance on a larger, more computationally expensive problem. With this in mind, a training set is constructed from the LOW selections and objective function values of the TRN problem. Specifically, the 10 subsequences of lengths two and three, with the largest γ -ratios are selected from the LOW set for TRN, and this training set is used to offline train SSHH. The offline trained hyper-heuristic, denoted T-SSHH-L-TRN, uses the same switching mechanism as T-SSHH-L described in the previous section. As all the training subsequences are drawn from one problem and evaluated on the remaining 11 problems, the leave-one-out cross validation methodology is not required. The results for the TRN problem are included for completeness.

The offline trained hyper-heuristic T-SSHH-L-TRN is executed 40 times on each of the 12 problems in the WDN domain. Table 13 contains the mean final objective function values for each problem in the WDN domain. The results demonstrate that T-SSHH-L-TRN outperforms SSHH on every problem (see Table 9), and outperforms T-SSHH-L on five out of 11 problems.

The overall optimisation results shown in Table 14 indicate that T-SSHH-L-TRN outperforms SSHH, and is slightly less effective than T-SSHH-L.

The overall differences in final log returns are tested for statistical significance. The results, shown in Table 15, indicate that the offline trained hyper-heuristic T-SSHH-L-TRN outperforms the SSHH hyper-heuristic on 11 of the WDN problems with 99% confidence.

Table 14: The mean final log return $\bar{\alpha}_f$, the mean percentage change, the mean number of iterations to a minimum, and the mean number of objective function evaluations.

	$\bar{\alpha}_f$	%	Min. Sel.	Obj. Eval.
SSHH	-2.9104	-88.7454	32916.5854	32238.0979
T-SSHH-L	-2.9248	-88.8740	34693.5292	30437.8688
T-SSHH-L-TRN	-2.9233	-88.8643	33901.3896	29479.8042

Table 15: The sample median difference \hat{d} , the sample median absolute deviation MAD, the sample mean difference \bar{d} , the standard deviation SD, the p -value, and the 99% confidence interval for $\bar{\alpha}_f(\text{T-SSHH-L-TRN}) - \bar{\alpha}_f(\text{SSHH})$ (excluding the TRN problem). Statistically significant results are shown in boldface.

\hat{d}	MAD	\bar{d}	SD	p -value	Conf. Int.
-0.0039	0.0023	-0.0138	0.1977	0.0000	$[-\infty, -0.0022]$

The result that subsequences drawn from TRN, which is the smallest problem in the WDN domain, can be used to improve the optimisation of EXN which is the largest, is notable. However, the improvement in performance is perhaps less surprising when one notes that the heuristic orderings in the LOW sets are very similar for TRN and EXN (see Table 8).

These experimental results demonstrate that it is possible to learn offline from a small, computationally inexpensive problem, and then use this knowledge to improve optimisation performance on larger, more computationally expensive WDN problems with the same objective function.

5 Conclusions

The sequence-based selection hyper-heuristic SSHH is able to produce viable solutions for the 12 WDN problems that are comparable in monetary cost and resilience to the cheapest solutions presented in Wang et al. (2015) but using significantly less computational resources. In addition, the two-point crossover heuristic is shown to perform well when compared with the five low-level heuristics employed in Kheiri et al. (2015).

The selection hyper-heuristic DBGen is used to generate an offline learning database of low-level heuristic selections and their objective function values across the 12 problems in the WDN domain. By employing the framework presented in Yates and Keedwell (2019), low-level heuristics and subsequences of heuristics can be identified in the database as being either effective or disruptive. Effective subsequences tend to decrease the objective function value, while disruptive subsequences tend to increase the objective function value. The most effective heuristic subsequences are used to offline train the HMM of the SSHH hyper-heuristic using the Baum-Welch learning algorithm. However, the Baum-Welch algorithm is unable to learn an effective optimisation strategy because the performance of the effective subsequences varies considerably during the optimisation process, and this variation can be quantified by the Spearman’s footrule metric. In order to test this hypothesis, subsequences that are

Downloaded from http://direct.mit.edu/evco/article-pdf/29/2/187/1921046/evco_a_00277.pdf by guest on 11 August 2021

effective when the objective function value is relatively low are selected from the database. These subsequences are used to train another HMM which is employed by the SSHH hyper-heuristic after the midpoint of the optimisation process. Although the method of switching between optimisation strategies is simple, it produces a small, but notable improvement in performance. The final experiment demonstrates that it is possible to learn offline from a small WDN problem which is computationally inexpensive, and transfer this learning to larger, more computationally expensive problems. Furthermore, this improvement in optimisation performance is statistically significant, with 99% confidence.

Although the offline learning gains are modest, it should be possible to improve on these results. For example, one obvious enhancement would be to improve the mechanism employed to switch between the optimisation strategies encoded in the initial and offline trained HMMs.

Another limitation of this study is that SSHH and the MOEAs are compared at only a single point; the cheapest viable solution. This weakness could be addressed by employing the multiobjective version of SSHH described in Walker and Keedwell (2016). This would allow SSHH to generate a Pareto front of solutions during optimisation, enabling a direct comparison with the Pareto fronts presented in Wang et al. (2015), and facilitating an analysis of the trade-offs between the conflicting design objectives of cost and resilience.

It is clear from the EXN* result in Section 4.1 that optimisation can also be improved by “tuning” the objective function constants a , b , and c for each problem. However preliminary experiments suggest that using the same objective function facilitates the generalisation exhibited by T-SSHH-L and the transfer learning exhibited by T-SSHH-L-TRN. A fuller investigation into learning when these constants vary for each problem is left to future research.

References

- Alperovits, E., and Shamir, U. (1977). Design of optimal water distribution systems. *Water Resources Research*, 13(6):885–900.
- Baños, R., Reca, J., Martínez, J., Gil, C., and Márquez, A. (2011). Resilience indexes for water distribution network design: A performance analysis under demand uncertainty. *Water Resources Management*, 25(10):2351–2366.
- Bishop, C. M. (2006). *Pattern recognition and machine learning*. New York: Springer.
- Blum, C., and Roli, A. (2003). Metaheuristics in combinatorial optimization: Overview and conceptual comparison. *ACM Computing Surveys*, 35:268–308.
- Bragalli, C., D’Ambrosio, C., Lee, J., Lodi, A., and Toth, P. (2008). *An minlp solution method for a water network problem*. Technical Report RC24495. IBM Research.
- Burke, E. K., Hyde, M. R., Kendall, G., Ochoa, G., Özcan, E., and Woodward, J. R. (2019). *A classification of hyper-heuristic approaches: Revisited*, pp. 453–477. Cham: Springer International Publishing.
- Burke, E. K., Hyde, M. R., Kendall, G., and Woodward, J. R. (2007). The scalability of evolved on line bin packing heuristics. In *IEEE Congress on Evolutionary Computation*, pp. 2530–2537.
- Cunha, M., and Sousa, J. (1999). Water distribution network design optimization: Simulated annealing approach. *Journal of Water Resources Planning and Management*, 125(4):215–221.

- Demšar, J. (2006). Statistical comparisons of classifiers over multiple data sets. *The Journal of Machine Learning Research*, 7:1–30.
- Derrac, J., García, S., Molina, D., and Herrera, F. (2011). A practical tutorial on the use of nonparametric statistical tests as a methodology for comparing evolutionary and swarm intelligence algorithms. *Swarm and Evolutionary Computation*, 1(1):3–18.
- Diaconis, P., and Graham, R. L. (1977). Spearman's footrule as a measure of disarray. *Journal of the Royal Statistical Society. Series B*, 39(2):262–268.
- Dietterich, T. G. (1998). Approximate statistical tests for comparing supervised classification learning algorithms. *Neural Computation*, 10:1895–1923.
- Drake, J. H., Kheiri, A., Özcan, E., and Burke, E. K. (2019). Recent advances in selection hyper-heuristics. *European Journal of Operational Research*, 285:143–165.
- Drake, J. H., Özcan, E., and Burke, E. K. (2015). A comparison of crossover control mechanisms within single-point selection hyper-heuristics using HyFlex. In *IEEE Congress on Evolutionary Computation*, pp. 3397–3403.
- Eusuff, M. A., and Lansey, K. E. (2003). Optimization of water distribution network design using the shuffled frog leaping algorithm. *Journal of Water Resources Planning and Management*, 129(3):210–225.
- Ezzeldin, R., Djebdjan, B., and Saafan, T. (2014). Integer discrete particle swarm optimization of water distribution networks. *Journal of Pipeline Systems Engineering and Practice*, 5(1):04013013.
- Farmani, R., Savic, D. A., and Walters, G. A. (2004). EXNET benchmark problem for multi-objective optimization of large water systems. In *Modelling and Control for Participatory Planning and Managing Water Systems, IFAC Workshop, Venice, Italy*.
- Fujiwara, O., and Khang, D. B. (1990). A two-phase decomposition method for optimal design of looped water distribution networks. *Water Resources Research*, 26(4):539–549.
- Geem, Z. W. (2006). Optimal cost design of water distribution networks using harmony search. *Engineering Optimization*, 38(3):259–277.
- Gessler, J. (1985). Pipe network optimization by enumeration. In *Speciality Conference on Computer Applications / Water Resources*, pp. 572–581.
- Hodges, J. L., and Lehmann, E. L. (1963). Estimates of location based on rank tests. *The Annals of Mathematical Statistics*, 34(2):598–611.
- Kheiri, A., and Keedwell, E. C. (2015). A sequence-based selection hyper-heuristic utilising a hidden Markov model. In *Proceedings of the Genetic and Evolutionary Computation (GECCO)*, pp. 417–424.
- Kheiri, A., and Keedwell, E. C. (2017). A hidden Markov model approach to the problem of heuristic selection in hyper-heuristics with a case study in high school timetabling problems. *Evolutionary Computation*, 25(3):473–501.
- Kheiri, A., Keedwell, E. C., Gibson, M. J., and Savic, D. (2015). Sequence analysis-based hyper-heuristics for water distribution network optimisation. *Procedia Engineering*, 119:1269–1277.
- Kim, J. H., Kim, T. G., Kim, J. H., and Yoon, Y. N. (1994). A study on the pipe network system design using non-linear programming. *Journal of Korea Water Resources Association*, 27(4):59–67.
- Lee, S. C., and Lee, S. I. (2001). Genetic algorithms for optimal augmentation of water distribution networks. *Journal of Korea Water Resources Association*, 34(5):567–575.

- Mohan, S., and Babu, K. S. J. (2010). Optimal water distribution network design with honey-bee mating optimization. *Journal of Computing in Civil Engineering*, 24(1):117–126.
- Raad, D. N., Sinske, A. N., and van Vuuren, J. H. (2010). Comparison of four reliability surrogate measures for water distribution systems design. *Water Resources Research*, 46(5):W05524.
- Rabiner, L. R. (1989). A tutorial on hidden Markov models and selected applications in speech recognition. *Proceedings of the IEEE*, 77(2):257–286.
- Reca, J., and Martínez, J. (2006). Genetic algorithms for the design of looped irrigation water distribution networks. *Water Resources Research*, 42(5):W05416.
- Rossman, L. A. (2000). *EPANET2 users manual*. U.S. Environment Protection Agency.
- Schaake, J. C., and Lai, F. H. (1969). *Linear programming and dynamic programming application to water distribution network design*. Technical report. M.I.T. Hydrodynamics Laboratory.
- Sherali, H. D., Subramanian, S., and Loganathan, G. V. (2001). Effective relaxations and partitioning schemes for solving water distribution network design problems to global optimality. *Journal of Global Optimization*, 19(1):1–26.
- Shirzad, A. (2017). Shortening the search time in optimization of water distribution networks. *Urban Water Journal*, 14(10):1038–1044.
- Veek, N., Repinek, M., and Mernik, M. (2017). On the influence of the number of algorithms, problems, and independent runs in the comparison of evolutionary algorithms. *Applied Soft Computing*, 54(C):23–45.
- Walker, D. J., and Keedwell, E. C. (2016). Multi-objective optimisation with a sequence-based selection hyper-heuristic. In *Proceedings of the 2016 Genetic and Evolutionary Computation Conference Companion (GECCO) Companion*, pp. 81–82.
- Wang, Q., Guidolin, M., Savic, D., and Kapelan, Z. (2015). Two-objective design of benchmark problems of a water distribution system via MOEAs: Towards the best-known approximation of the true Pareto front. *Journal of Water Resources Planning and Management*, 141(3):04014060.
- Yates, W. B., and Keedwell, E. C. (2019). An analysis of heuristic subsequences for offline hyper-heuristic learning. *Journal of Heuristics*, 25(3):399–430.
- Zheng, F., Zecchin, A. C., and Simpson, A. R. (2013). Self-adaptive differential evolution algorithm applied to water distribution system optimization. *Journal of Computing in Civil Engineering*, 27(2):148–158.

Appendix: The Wilcoxon Test

The Wilcoxon test assumes that the paired values of $\alpha_f(A)$ and $\alpha_f(B)$ are independently drawn. The use of cross-validation to determine generalisation violates this assumption as the training sets overlap. This overlap may prevent the statistical test from obtaining a good estimate of the amount of variation that would be observed if the training sets were completely independent. As a consequence, when cross-validation is used, the results of the statistical tests must be viewed as approximate rather than rigorously correct (see Dietterich, 1998).

The null hypotheses of the Wilcoxon test is that the median difference between pairs of observations is zero. As the differences between the $\alpha_f(A)$ and $\alpha_f(B)$ values are not symmetrically distributed around the median, the Hodges–Lehmann estimate of the median is used instead (see Hodges and Lehmann, 1963).

All statistical tests were performed using the R language for statistical computing.